Untitled

by anay

General metrics

31,911	5,353	516	21 min 24 sec	41 min 10 sec
characters	words	sentences	reading time	speaking time

Score

This text scores better than 49% of all texts checked by Grammarly

Writing Issues

375	230	145
Issues left	Critical	Advanced

Plagiarism

This text seems 100% original. Grammarly found no matching text on the Internet or in ProQuest's databases.

Writing Issues

231	**Correctness**
1	Misplaced words or phrases
28	Misspelled words
6	Incorrect verb forms
7	Mixed dialects of english
1	Text inconsistencies
59	Incorrect punctuation
6	Pronoun use
8	Incorrect noun number
22	Confused words
12	Faulty subject-verb agreement
34	Determiner use (a/an/the/this, etc.)
7	Conjunction use
3	Misuse of modifiers
14	Incorrect phrasing
12	Wrong or missing prepositions
6	Ungrammatical sentence
4	Improper formatting
1	Closing punctuation

139	**Clarity**
3	Hard-to-read text
7	Paragraph can be improved
9	Wordy sentences
22	Intricate text
98	Passive voice misuse

5 **Delivery**

5 Tone suggestions

Unique Words

Measures vocabulary diversity by calculating the percentage of words used only once in your document

20%

unique words

Rare Words

Measures depth of vocabulary by identifying words that are not among the 5,000 most common English words.

43%

rare words

Word Length

Measures average word length

4.6

characters per word

Sentence Length

Measures average sentence length

10.4

words per sentence

Untitled

Abstract—In recent years, facial research recognition[1] has advanced to a point where it can reliably detect faces in controlled settings, but it[2] still struggles with occlusion and in real world[3] scenarios. This paper reviews different models proposed to deal with occlusion. This paper also comments on the various strategies to[4] used to detect occlusion.

For[5] this paper, I have analyzed[5,6,7] the architectures proposed.[5] Along side[8] that, there is a short analysis of the common datasets that are used for training and testing these models.

This paper found that

1) Models that combine both global and local features out-perform those who do not.

2) No model attained a satisfactory score when tested on real occluded datasets.

3) Several models are constrained by their reliance on facial landmark detection models/algorithms.

Index Terms—Computer Vision, Facial Expression Recognition[9], Occlusion Handling, Deep Learning

INTRODUCTION

FACIAL Expression Recognition, or FER[10] has undergone numerous advancements in recent years. We are now at a point where we have models that can perform FER with high degrees of accuracy in laboratory settings [1].

However, occlusion still poses a problem to reliable FER. Occlusion refers to the phenomenon which occurs when a portion of a 3D scene is not visible from the camera or another sensor [2]. Many real world applications will need to deal with occlusion in order to be of any use. Keeping this in mind, it is essential that researchers deal with the occlusion problem sufficiently.

Researchers have been trying to deal with occlusion in FER (OFER) for over 20 year now, with Bourel et al. publishing one of the earliest works on the topic [3]. Since then numerous models have been proposed to deal with the issue. There have been several reviews on methods to detect FER have been published and among the first to include methods that deal with occlusion was Bettadapura [4] in 2012. A much more comprehensive survey was conducted by Zhang et al. [5] in 2018.

BACKGROUND

In this section, we shall first discuss the standard FER pipeline. This emerged over decades of research related to facial expressions primarily by psychologists. In the 1990s and 2000s, when machine learning techniques were first applied, this pipeline solidified [14][15]. The pipeline is given as:

1) Face Detection

2) Preprocessing

3) Feature Extraction

4) Classification

A. Occlusion Scenarios

There are numerous scenarios in real world[23] settings where the face is subject to varying degrees of occlusion. Face masks, glasses and caps [12][18] all can cause occlusion and lead to loss of effectiveness in FER models.

Pose can also contribute to occlusion. When someone turns their head, for example, several import[24] features of the face get hidden,[25] again causing the[25] effectiveness to be hit [19].

Another scenario that leads to occlusion is termed ”self-occlusion”, where the person blocks an area of their face, either deliberately or subconsciously. For example, someone rubbing their eyes [19].

B. Why Occlusion Hurts FER

If pixels are simply[26] missing, then that is information loss. This[27] is very bad as it is, but occlusion is a step further. It causes large scale[28] spatial changes in the face [20].

DATASETS

A major challenge researchers face in this field is the avail-[29] ability of datasets. There exist several datasets of human faces. These can be categorized[7,30,31] as either in-the-lab datasets or in-the-wild datasets. Examples of in-the-lab datasets include CK+ [6], MMI [7], Oulu-CASIA [8] and JAFFE [16]. Examples of in-the-wild datasets include RAF-DB [9], AffectNet [10], SFEW [11] and FER2013 [17].

One thing to note about all these datasets is that none of them contain[32] images with partial occlusion. All of them are clear face images. To the best of my knowledge, there is only one dataset that has collected faces with organic occlusions,

which is the FED-RO dataset [12].

RAF-DB contains 30,000 facial images annotated with basic or compound expressions by 40 trained human coders. The Extended[33] Cohn-Kanade dataset (CK+) contains 593 video sequences recorded from 123 subjects. MMI[34] dataset includes more than 30 subjects of both genders. There are 79 sequences[35] of each subject. Each begins and ends with neural[36] facial expression. Oulu-CASIA dataset contains six prototypic expressions from 80 people between 23 to[37] 58 years old [12].

SFEW The Static Facial Expressions in the Wild (SFEW) dataset[38] is created by selecting static frames from Acted[39] Facial Expressions in the Wild (AFEW) [13] This dataset contains 95 subjects in total and covers unconstrained facial expressions, different head poses, ages, and occlusions. In total[40] there are 663 well-labelled usable images. The Japanese Female Facial Expression[41] (JAFFE) database contains 213 samples of posed expressions from 10 Japanese females. Each person has 3~4 images with each of six basic facial expressions and one image with a neutral expression.

MODELS REVIEWED

This paper has reviewed the below[42] 9 models:

1) Patch-Gated CNN (PG-CNN) [21]

2) Patch-Based ACNN (pACNN) [12]

3) Global-Local-Based ACNN (gACNN) [12]

4) OADN [22]

5) WGAN for OFER [23]

6) WLS-RF [24]

7) Privileged Information CNN (PI-CNN)[25]

8) RGBT [26]

9) CT-DBN [27]

Out of these, PG-CNN [21], pACNN [12], gACNN [12], OADN [22] and WLS-RF [24] can be classified[43] as models that break[44] down the face into patches, determine[45] how useful each patch is and[46] then[47] use this information for classification. WGAN for OFER [23] and PI-CNN [25] uses[48] a generative approach to tackle occlusion. They deal with occlusion by trying to recreate the face if there was no occlusion in the first place.

A. Patch-Gated CNN and Patch-Based ACNN

PG-CNN is based[49] on the first 9 convolutional layers of VGG-16 [30]. This[50] gives feature maps of dimensions 512x28x28 .The feature maps from VGG-16 is[51] then fed[52] through a Region Decomposition algorithm which was proposed in [29] where 68 land mark[53] points are found.[54] Out of these, 16 are selected[55] to cover eyebrow corners, eye corners, nose bridge, tip of the nose, corners[56] of the mouth. 8 more points were derived. 4 midpoints for the eyes and eyebrows were calculated[57] and[58] 4 midpoints were[59] calculated[60] of the cheeks.[61] The 512x28x28 feature map is then cropped[62] around these 24 landmarks to give 24 local feature maps of dimension 512x6x6. Then, each of these 24 local feature maps are[63] processed independently in their[64] own patch-gate units. In each PGU, the feature maps go through 2 convolutions keeping[65] the dimensions same.[66] This output is then split[67] along 2 branches.

The first branch does feature encoding. It takes the feature map and encodes it into a 64-dimensional vector (512x6x6 to 128x3x3 to 64) [21].

The second branch is what the authors call an “attention net”, which gives how occluded that patch is. The attention net gives an output between 0 and 1 via a sigmoid and represents unobstructed-ness.[68] The output of the attention net is a simple scalar. The final output of the PGU (Patch Gate Unit) is just the occlusion score times the feature vector. The weighted feature vectors from all 24 PGUs are then concatenated[69] into a single large vector which[70] is passed[71] through 3 fully connected layers and finally a softmax classifier that outputs the probability for 7 different emotions, namely angry,[72] disgust, fear, happy,[73] sad,[74] surprise, neutral.[75] The architecture is represented[76] in Figure 1 [21].

It is worth noting that this model was trained[77] on synthetic occlusions. The occlusions added were constrained to the size range of 96x96 to 128x128 pixels. This range was chosen[78] because smaller occlusions had little effect on the model [21].

Patch-Based ACNN has an architecture identical to PG-CNN with[79] the sole difference being that PG-CNN has 3 fully connected layers and pACNN has 2. The similarity can be explained[80] simply by realising[7] that the authors of both [21] and [12] are the same and[81] pACNN is a more refined version of PG-CNN [12].

B. Global-Local-Based ACNN

gACNN, as proposed in [12] is[82] a further refinement of the

PG-CNN architecture. While the PG-CNN [21] and pACNN [12] architectures are good at locating occlusions, what they sometimes miss is the larger picture. It is now enough to know that this part of the face has untrustworthy information, we also need to know the greater details of the face to place the untrustworthy patch in a broader context [12].

To solve this, the authors took the identical framework of the pACNN, but added a global gated unit (GGU) which is identical in function and structure to the PGU, but this works on the whole face. There is only one other notable difference between gACNN and pACNN, which is that instead of using the first 9 convolutional layers of VGG-16, they use the first 12, this is done to get slightly larger spatial dimensions for the feature map that they send to the GGU. Figure 2 shows the architecture of gACNN [12].

C. Occlusion Aware Deep Network (OADN)

This model uses ResNet-50 as its backbone, instead of VGG-16. It uses all the convolutional layers of ResNet50 [31], the only changes it does is removing the average pooling and fully connected layers and changing the stride of conv4 1 from 2 to 1 to get a larger feature map. This is because both branches need as detailed spatial information that they can get [22] .

After ResNet, The model splits into 2 branches. The first branch is the Landmark-guided Attention Branch or LAB. Firstly, a pretrained landmark detector, SAN [32] which is trained on the 300W dataset [33] extracts 68 facial landmarks.

It outputs both the coordinates and a confidence score between 0 to 1, basically saying how confident the SAN is in its detection. From these 68, 24 points are selected. 16 are directly selected for covering eyes, eyebrows, nose and mouth, 8 are recomputed for cheeks and the areas between eyes and eyebrows. 2 things are important to note here. First, there seems to be some discrepancy in how many points are taken from the 68. Though they first claim that 16 are directly selected and 8 are recomputed, in a later section they claim that 18 are directly selected. Secondly, we are not told how the recomputed points are exactly recomputed. Though, considering how they cite the gACNN papers [12], it is likely that they use the very same method, but it is not explicitly mentioned. For the recomputed points, the confidence score is assigned of the parent point with the lower score. Say a point is recomputed from points with confidence scores of 0.76 and 0.59, the recomputed point will have a score of 0.59 [22]. The model applies a confidence threshold (denoted by T) which the authors set to 0.6. Any point with confidence score below T is considered to be occluded and hence discarded. Now for the interest points which are not occluded, a 2D gaussian heat map is applied to the point where the landmark is situated. Think of the landmark as a bulb. The landmark will have the brightest light, and the brightness of the light will fade as you go farther away from the bulb. The gaussian heat map tells how bright a point will be. Mathematically, it is expressed as:

G(x, y) = exp

□

− (x − x0)2 + (y − y0)2

2σ2

□

Where σ controls how wide the glow spreads. For occluded points, the heat map produced has all zeros [22].

The heatmaps are then downsampled[115] to be of 14x14 resolu-[116] tion using linear interpolation, so as to[117] match the resolution of the backbone's feature map. Linear interpolation is used here instead of any other method because it compresses the heatmap while preserving the general shape of the gaussian.[118] The backbone produces a feature map of dimensions 2048x14x14, so the heat map applied[119] is applied[120] across all 2048 channels. Now, we have 24 such heatmaps, all of which are applied on the same feature map to get 24 separate copies, all masked differently. Each of these 2048x14x14 feature maps is then global average pooled[121] into a 2048D vector. All of these vectors represent what the network "sees" around a landmark. Then these 24 vectors are max pooled[122] across the 24 landmarks, that is, for each of the 2048 dimensions, the maximum value across the 24 vectors is taken,[123] fusing the information from all the regions into a single 2048D vector. This 2048D vector is then reduced to a 256D vector using a fully connected layer and is[124] then fed[125] into a softmax classifier which[126] gives the probability distribution over the 7 classes, namely angry, disgust, fear, happy, sad, surprise, neutral. This branch is trained according

to cross-entropy loss [22].

The second part of the OADN is the facial region branch. This branch takes the same 2048x14x14 feature map and divides the 14x14 spatial grid into K non-overlapping blocks. The authors found that K = 4 gives the best results. Each block, independently, is reduced to a 2048D vector, which is passed[127] through its own FC layer, which reduces it to a 256D vector, which is passed on[128] to a softmax layer which[129] predicts the expression category. Each block's classifier is trained[130] independently using cross entropy[131] loss and[132] the total loss is the sum of individual losses across all 4 regions [22].

There is some ambiguity as to how the model actually predicts a single class. The authors have told[133] that the final output is the average of the LAB output and the FRB output. The ambiguity arises on[134] how this operation is exactly done[135] ?[134]

Is it

LAB + average of 4 regions

2

Or,

LAB + region1 + region2 + region3 + region4

5

It is likely that the first interpretation is[136] used[137] since the paper itself admits that the LAB is more effective on its own than the FRB on its own [22].

OADN uses 2 branches because the LAB is only good[138] as the points it selects. So it is plausible that the image has a person wearing a mask, making the landmark detection algorithm to[139]

perform poorly, leading to the LAB performing poorly. The LAB is the main heavyweight in the OADN,[140] the FRB plays a supporting role, stepping up when the LAB is ineffective. The greatest weakness of the OADN is its dependence on the SAN. The two branches are trained[141] according to a combined loss function given by

L = λLLAB + (1 − λ)LF RB

where[142] λ determines the relative importance of each branch. An important distinction between OADN and the gACNN family[143] is that the gACNN family learns to detect which parts are occluded, while the OADN relies completely on the SAN's assessment on[144] occlusion. The largest weakness of OADN is its dependence on SAN [22].

The architecture of OADN is shown[145] in figure[146] 3

D. WGAN for OFER

At the very start, all the images are preprocessed[147] using MTCNN [34]. This network detects the face, the centres of the eyes and aligns the face so that the eye centres are horizontal. The face region is cropped[148] and resized to be of the[149] size 128x128 and[150] the values of the pixel are normalized[7,151,152] to 0 to 1 [23] .[153]

This model proposes a variation of WGAN [35] which[154] itself is a variation of the GAN network [36]. The model consists of 3 networks. A generator and 2 discriminator.[155] The generator is a convolutional autoencoder [38] which takes input and output of the size 3x128x128. The input is an occluded image and[156] the output is an reconstructed[157] image

without occlusion. A huge shortcoming of WGAN for OFER is that it is trained[158] on synthetic occlusions. In the encoder, the depth of the image increases until it is eventually 512 layers deep. Each convolution uses a 3x3 kernels[159] and has stride[160] 1.[161] Every convolution is followed[162] by batch normalization[7,163] and ReLU [37]. The output of the decoder[164] is 512x8x8. The bottleneck consists of 2 fully connected layers with 1024 neurons where the incoming feature map is flattened and compressed into a 1024D vector. The decoder has a similar structure as[165] the encoder but[166] in the reverse direction and uses transpose convolutions [23] .[167]

Both of the discriminators are based on VGG-19 and have the same architecture. The only real difference between them is in their fully connected layers. Both use 3x3 kernels, stride[168] 1 and ReLU activation. After every convolution, the feature maps are downsized by pooling. Though it is not mentioned[169] which type of pooling is used,[170] it is very likely that it is[171] max pooling considering[172] that VGG-19 uses it. The pooling layers uses[173] 2x2 filters with stride[174] 2.[175] After all this, both the models[176] share 3 fully connected layers. First 2 of those contain 4096 neurons and the third one contains 1000 neurons [23].

Both of the discriminators are based on VGG-19 and have the same architecture. The only real difference between them is in their fully connected layers. Both use 3x3 kernels, stride[177] 1 and ReLU activation. After every convolution, the feature maps are downsized by pooling. Though it is not mentioned[178] which type of pooling is used,[179] it is very likely that it is[180] max

pooling considering[181] that VGG-19 uses it. The pooling layers uses[182] 2x2 filters with stride[183] 2.[184] After all this, both the models[185] share 3 fully connected layers. First 2 of those contain 4096 neurons and the third one contains 1000 neurons [23].

Now coming[186] to the difference between the generators. D1 is the real/fake discriminator. After the FC layers, a sigmoid function gives an output which tells[187] how real or fake it is. An important thing to note here is that if the model is using Wasserstein distance, it shouldn't really be using a sigmoid function. D1's objective is to push the generator to make more realistic non-occluded images [23].

D2 is the expression classifier which[188] after the FC layers[189] sends[190] the input through a softmax function and outputs probabilities over 7 expressions: happiness, sadness, surprise, disgust, fear, anger, neutral [23].

Both discriminators are trained[191] on 3 types of inputs, namely generated de-occluded images, original un-occluded images and irrelevant-area occluded images. Irrelevant area occluded images are images where the occlusion is on areas which do[192] not carry a lot of information relating to facial expressions. This[193] is done[194] so that D2 learns authentic expression features since reconstructed images will carry subtle differences from natural images [23].

The model is trained[195] using 4 loss functions. The first is weighted[196] reconstruction loss. This[197] measures[198] pixel level similarity[199] between generated and original non-occluded images. The second loss function is triplet loss. To calculate this,

we define 3 terms. d1 which[200] is the difference between the generated and original image in the occluded area and[201] d2 which[202] is the difference between the generated and occluded images in the occluded area.[203] The objective is to make d1 small and d2 large. Another factor β is introduced[204] to make sure that d2 is not marginally larger than d1 .[205] The third loss function is adversarial[206] loss from the WGAN. The final loss function is the classification loss which[207] is the standard cross entropy[208] loss on D2 [23].

The discriminator and generator are trained[209] alternatively .[210]

For every time the discriminators are trained, the generator is trained[211] 5 times.

The architecture of WGAN for OFER is given[212] in figure[213] 4

E. WLS-RF

Feature extraction for this model is done using SDM, or supervised descent method [39]. The model uses 3 features in total, of which 2 are derived[214] from the landmarks detected by SDM, the third represents the edges, which in this model are detected by histogram[215] of oriented gradients [40].

The model then trains random forests on the image, but not the entire image. A triangular mesh is defined[216] using the facial landmarks and[217] a mask is generated by randomly selecting connected triangles. This mask is allowed to grow until it covers 1

10

th of[218] the original image [24].

The model detects occlusion using a hierarchical autoen-[219]

coder[220] which assigns[220] a confidence score to each triangle.

The autoencoder is trained on non-occluded data however,[221] with 25% masking noise. The final prediction is simply the prediction of each triangle weighted by its confidence [24].

This model was also shown[222] to be capable of detecting facial action units [41].

The architecture of WLS-RF is given[223] in figure[224] 5.

F. PI-CNN

This model deal[225] with occlusion by treating the non-occluded images as privileged information [42]. The model consists of 2 CNNs, one classifies occluded facial expressions, this is[226] our primary classifier and[227] the other classifies non-occluded facial expressions. Both are pre-trained using cross entropy[228] loss, after which the weights of the non-occluded network are fixed.[229] Now, the occluded network is fine tuned[230,231] with guidance in label space and feature space [25].

For label space, the two CNNs are subject to two constraints, the first is a similarity constraint which ensures that the output[232] of both aren't wildly different and[233] the second is loss[234] inequality regularization[7,235] [43]. The intuition behind it that[236] the privileged information will always be more descriptive than than the non-privileged information, hence the output of the privileged[237] network must always be closer to the base truth [25].

For guidance in feature space, a GAN-like network is introduced to ensure that the occluded network learns robust patterns that help in classification. The GAN-like network

is only used in training and is only used to ensure that the classifier[238] learns the desired kind of information, it[239] is not used[240] during inference [25].

The architecture of PI-CNN is shown[241] in figure[242] 6

G. RGBT

The first step in this model is detecting the face using the Viola-Jones face detector [44]. Then a gabor[243] filter [46] bank is applied giving[244] us a series of gabor[245] images. Then, templates are extracted[246] using a Monte Carlo algorithm [45].

Then an image from a certain class is selected and[247] a random position is chosen,[248] on this, one of the scales from the gabor[249] filter bank is applied randomly. This[250] is done[251] to[252] 4 different sizes of templates. This mechanism is introduced[253] to deal with occlusion. Intuitively, occlusion will only cover a small area of the face, so if you randomly sample many such filters, you are bound to cover the entire face, when there is and isn't occlusion [26].

The next step is template matching,[254] For the image, every template is applied, but not directly. It searches within a region twice the width and height of the template centered[255] on the original template's position. Within this area, the template is slid[256] across all positions and the distance is computed using one of two sparse metrics. By this point, we have a 28000 dimensional[257] feature vector for each image, which will eventually have to be reduced [26].

We now come to training our classifier. A one against all SVM [47] is trained[258] for each emotion and[259] then pruning starts

to reduce dimensionality. The importance of each template is calculated from the weights of the SVM and in the first iteration of training, the bottom half of templates are removed. This is done until the number is less than 1000, after which the templates are reduced 50 per iteration until the number goes below 100, after which they are reduced 10 per iteration. We should expect to get between 10 to 20 final templates [26]. In final classification, the emotion corresponding to the emotion with the highest decision value is the prediction [26]. The framework of RGBT is given in figure 7.

H. CT-DBN

This model uses a dual branch approach to deal with occlusion as seen earlier. Here we have a CNN and a Swin Transformer [48] which is a variation of the Transformer architecture [49].

First, the raw image goes through face detection using the Dlib library. Data augmentation is applied during training where the image is randomly flipped horizontally and color jitter is applied. From this image, local binary pattern (LBP) is applied to encode local texture information [27].

At this point, the model splits into 2 branches, first is the local feature extraction branch (LFEB) and the other is the global feature extraction and feature fusion branch (GFE-FFB). Both branches simultaneously process their inputs [27].

In the LFEB branch, the LBP image is fed into the first 5 layers of ResNet-18. This is more or less a standard CNN which outputs a prediction P red1. This is not used for the

final classification, it is only to train the LFEB [27].

The GFE-FFB is a standard Swin Transformer along with all the accompanying embeddings and reshaping. The Swin Transformer consists of 2 sub-blocks, both are processed consecutively. The first sub-block is regular window multi-head self-attention (WMSA) and the second sub-block is shifted window multi-head self-attention (SWMSA) [27]

Then, global and local feature weights are computed, which are then used to compute a combined feature weight. This feature weight is used to calculate hybrid feature maps by using the global and local feature maps that were produced by both the branches [27].

The hybrid feature maps are passed through a fully connected layer to get P red2 which is our final prediction [27].

The architecture of CT-DBN is given in figure 8.

V. DETAILED ANALYSIS OF THE PROPOSED MODELS

Each of the models proposed is trained and tested on different datasets. Table 1 shows all the different datasets that the models used as either testing or validation. Furthermore, a lot of these models train on custom datasets constructed from images from two or more of the above datasets. This is the greatest limitation of this review. Due to constraints on time and competence, these models could not be tested on a level playing field. We will be relying completely on the figures provided by the original authors.

A factor that must be noted is that almost all of the above datasets were originally created as simple, non-occluded

datasets. Occlusion-FERPlus contains real occlusions. It was collected by reviewing images in FERPlus and selecting those with genuine occlusions. Occlusion-FERPlus, as a result[305] contains a relatively small dataset of 605 samples [28].

The only dataset that was collected[306] from the start as a[307] occlusion dataset is FED-RO [12]. This[308] was collected[309] by mining Bing and Google search with search queries like ”smile+face+glasses”. Then human annotators labelled the images. Then, images that overlapped with RAF-DB and AffectNet were removed.[310] This[311] too[312] is[313] a relatively small dataset of around 400 samples. The test results on this dataset are perhaps the closest to real world[314] results we will get. This[315] shows that there is still large[316] scope for improvement since all the models who[317] tested on FED-RO reported sub-75% accuracy.

For purposes of our comparison, we will be using the figures of the models for the RAF-DB [9] and CK+ [6] datasets. This[318] is not a perfect way to compare as[319] we are still comparing on two different datasets, but using[320] the models[321] performance relative to each other, we can get a rough idea of how the different methods fare.

Below is the accuracy of various models while testing with RAF-DB dataset

Note the absence of RGBT [26] in both of the above tables. This[322] is because there is no single figure provided by the authors.[323] They provided results for their model only on specific types and sizes of occlusion.

Another important thing to keep in mind is that unless

specified otherwise, most of the above models were trained[324] on a combination of several datasets, often times[325,326] including RAF-DB and CK+. Both[327] of which were not[328] collected[329] with occlusion in mind.

This issue gets highlighted when models are tested[330] on datasets like FED-RO [12] and Occlusion-FERPlus [28]. Table 4[331] shows the accuracy of the models that were tested[332] on the occlusion specific[333] datasets

Out of all the models compared, we can see that the model which tried to reconstruct the image, WGAN for OFER [23] and PI-CNN [25], performed worse on CK+ and RAF-DB datasets than models which[334] explicitly detected occlusion and dealt with it on a patch-wise basis.

It is certainly interesting to note that PI-CNN [25], a ”reconstructivist” model[335] scored higher accuracy than OADN [22], a patch based[336] model, even though it scored higher while testing on the RAF-DB dataset. This[337] suggests that privileged information could be more effective for real world[338] facial recognition. Unfortunately, I was only able to find one such model which[339] used this principle and[340] nothing can be said[341] definitively as of yet.

The high accuracy for CT-DBN [27] and OADN [22] suggest that multiple branched models that base their inference on both local and global features fare better than those who do not.

The case of WLS-RF [24] seems to contradict[344] the above statement, but it can be explained[342] by realizing[7,343] the fact that the

model trains multiple overlapping random forests. Though this does not directly learn global features, it makes sure that the model learns relations between regions that may be spatially ar[345] away from each other. Despite this, WLS-RF [24] does not consider global features. It learns the relationship between features that are far away from each other but[346] this is not the same as looking at the whole picture.

A final point to note is that CT-DBN [27] has no explicit occlusion handling mechanism. Its architecture is occlusion resistant to a large degree on its own.

VI. CONCLUSION

A. Summary of findings

1) This review can state that models that consider both local and global features fare better against occlusion than models that just[347] account for local features, at least on synthetic occlusion testing sets.

2) A massive weakness of models that rely on facial landmarks[348] is the landmark detection method they choose[349] to use. This[350] means that architectural decisions on landmark detection models/algorithms has[351] a downstream effect, sometimes limiting on[352] the facial expression recognition model. This[353] is a large topic and warrants a closer look in a separate review.

3) Based on the figures available, there is no model that performs satisfactorily on OFER in real world condi-[354] tions, as seen in the low accuracies scored by pACNN, gACNN, OADN and PI-CNN on the FED-RO dataset.

While models have gotten reasonably good at detecting synthetic occlusions.

B. Limitations

The largest and limitation[355] in this field is the unavailability of a large, high quality facial[356,357] images datasets that has occlusions[356,357]. There exists an opportunity for individual researchers to collect such a dataset. This[358] alone can help massively in OFER.

C. Future Directions of Interest

1) This is in no way a comprehensive review. There are several models, like those proposed in [28][52][53][54][55] have not been covered in this review. Further, this review has not compared models on a level playing field, which has hindered the analysis presented.

2) As stated prior,[359] it will be a huge contribution to collect a dataset consisting only of naturally occuring[360] occluded faces.

3) There is still a need to propose a model that can reliably detect facial expressions in real world[361] scenarios. No model scored a high accuracy on occlusion specific[362] datasets.

4) Through this review, we can identify that networks that account for both local and global features[363] fare better than those who do not. This[364] tells us that any future models, which aim to achieve high accuracy in OFER are likely to account for global and local features.

5) The relatively good performance of PI-CNN [25] on the FED-RO[365] [12] dataset shows us that models using

privileged information could potentially work better in real world[366] scenarios. Unfortunately, there are not many models using[367] privileged information to deal with[367] OFER. This[368] could be a future area of academic interest.

6) The high accuracy scored by WLS-RF [24] on the CK+[369] [6] dataset shows that the approach of training multiple[370] classifiers on overlapping regions could help in the creation[371] of a high accuracy[371] OFER model. Especially when[372] you consider that WLS-RF [24] does not account for global features.

7) The relatively high accuracy scored by CT-DBN [27] when it doesn't even have a specific occlusion-handling mechanism, shows that a building on this architecture, by adding some[373] occlusion handling mechanism could make it even better of[374] a classifier. But this depends on occlusion handling mechanisms to be good enough to give[375] it the necessary boost.

1.	recognition research	Misplaced words or phrases	Correctness
2.	~~, but it~~ → . However, it	Hard-to-read text	Clarity
3.	~~real world~~ → real-world	Misspelled words	Correctness
4.	~~to~~	Incorrect verb forms	Correctness
5.	*For this paper, I have analyzed the architectures proposed.*	Paragraph can be improved	Clarity
6.	~~analyzed~~ → analysed	Mixed dialects of English	Correctness
7.	*analyzed; categorized; realising; normalized; normalization; regularization; realizing*	Text inconsistencies	Correctness
8.	~~Along side~~ → Alongside	Misspelled words	Correctness
9.	~~Recogni~~ → Recognition	Misspelled words	Correctness
10.	FER,	Incorrect punctuation	Correctness
11.	~~which occurs~~ → that occurs	Pronoun use	Correctness
12.	~~real world~~ → real-world	Misspelled words	Correctness
13.	~~in order to~~ → to	Wordy sentences	Clarity
14.	researchers must deal	Wordy sentences	Clarity
15.	~~year~~ → years	Incorrect noun number	Correctness
16.	then,	Incorrect punctuation	Correctness
17.	that have	Pronoun use	Correctness
18.	, and	Incorrect punctuation	Correctness
19.	*This*	Intricate text	Clarity

20.	, primarily	Incorrect punctuation	Correctness
21.	*machine learning techniques were first applied*	Passive voice misuse	Clarity
22.	*is given*	Passive voice misuse	Clarity
23.	~~real world~~ → real-world	Misspelled words	Correctness
24.	~~import~~ → important	Confused words	Correctness
25.	*hidden, again causing the effectiveness to be hit [19].*	Paragraph can be improved	Clarity
26.		Tone suggestions	Delivery
27.	*This*	Intricate text	Clarity
28.	~~large scale~~ → large-scale	Misspelled words	Correctness
29.	~~avail~~ → availability	Misspelled words	Correctness
30.	*be categorized*	Passive voice misuse	Clarity
31.	~~categorized~~ → categorised	Mixed dialects of English	Correctness
32.	~~contain~~ → contains	Faulty subject-verb agreement	Correctness
33.	The Extended	Determiner use (a/an/the/this, etc.)	Correctness
34.	The MMI	Determiner use (a/an/the/this, etc.)	Correctness
35.	~~quences~~ → sequences	Misspelled words	Correctness
36.	a neural	Determiner use (a/an/the/this, etc.)	Correctness

37.	~~to~~ → and	Conjunction use	Correctness
38.	The dataset	Determiner use (a/an/the/this, etc.)	Correctness
39.	the Acted	Determiner use (a/an/the/this, etc.)	Correctness
40.	total,	Incorrect punctuation	Correctness
41.	The Expression	Determiner use (a/an/the/this, etc.)	Correctness
42.	~~below~~ → following	Misuse of modifiers	Correctness
43.	*be classified*	Passive voice misuse	Clarity
44.	~~break~~ → breaks	Faulty subject-verb agreement	Correctness
45.	~~determine~~ → determines	Faulty subject-verb agreement	Correctness
46.	~~is and~~ → ,	Paragraph can be improved	Clarity
47.	, and	Incorrect punctuation	Correctness
48.	~~uses~~ → use	Faulty subject-verb agreement	Correctness
49.	*is based*	Passive voice misuse	Clarity
50.	*This*	Intricate text	Clarity
51.	~~is~~ → are	Faulty subject-verb agreement	Correctness
52.	*is then fed*	Passive voice misuse	Clarity
53.	~~land mark~~ → landmark	Misspelled words	Correctness

54.	*are found*	Passive voice misuse	Clarity
55.	*are selected*	Passive voice misuse	Clarity
56.	and corners	Conjunction use	Correctness
57.	*4 midpoints for the eyes and eyebrows were calculated*	Passive voice misuse	Clarity
58.	, and	Incorrect punctuation	Correctness
59.	for the cheeks were	Incorrect phrasing	Correctness
60.	*4 midpoints were calculated*	Passive voice misuse	Clarity
61.	~~of the cheeks~~	Wrong or missing prepositions	Correctness
62.	*is then cropped*	Passive voice misuse	Clarity
63.	~~are~~ → is	Faulty subject-verb agreement	Correctness
64.	~~their~~ → its	Incorrect noun number	Correctness
65.	, keeping	Incorrect punctuation	Correctness
66.	the same	Determiner use (a/an/the/this, etc.)	Correctness
67.	*is then split*	Passive voice misuse	Clarity
68.	unobstructiveness	Incorrect phrasing	Correctness
69.	*are then concatenated*	Passive voice misuse	Clarity
70.	, which	Incorrect punctuation	Correctness
71.	*is passed*	Passive voice misuse	Clarity

72.	~~angry~~ → anger	Incorrect phrasing	Correctness
73.	~~happy~~ → happiness	Incorrect phrasing	Correctness
74.	~~sad~~ → sadness	Incorrect phrasing	Correctness
75.	and neutral	Conjunction use	Correctness
76.	*is represented*	Passive voice misuse	Clarity
77.	*this model was trained*	Passive voice misuse	Clarity
78.	*This range was chosen*	Passive voice misuse	Clarity
79.	, with	Incorrect punctuation	Correctness
80.	*The similarity can be explained*	Passive voice misuse	Clarity
81.	, and	Incorrect punctuation	Correctness
82.	],	Incorrect punctuation	Correctness
83.	*sometimes miss is the larger picture.*	Ungrammatical sentence	Correctness
84.	, which	Incorrect punctuation	Correctness
85.	~~the first~~ → The first	Confused words	Correctness
86.	~~the feature~~ → The feature	Confused words	Correctness
87.	The architecture	Confused words	Correctness
88.	~~is~~ → are	Faulty subject-verb agreement	Correctness
89.	*This*	Intricate text	Clarity
90.	~~that they~~ → as they	Conjunction use	Correctness

91.	~~The model~~ → the model	Confused words	Correctness
92.	The branch	Determiner use (a/an/the/this, etc.)	Correctness
93.	, which	Incorrect punctuation	Correctness
94.	*is trained*	Passive voice misuse	Clarity
95.	],	Incorrect punctuation	Correctness
96.	*are directly selected*	Passive voice misuse	Clarity
97.	*are recomputed*	Passive voice misuse	Clarity
98.	*are taken*	Passive voice misuse	Clarity
99.		Tone suggestions	Delivery
100.	section,	Incorrect punctuation	Correctness
101.	*are not told*	Passive voice misuse	Clarity
102.	*are exactly recomputed*	Passive voice misuse	Clarity
103.	Though,	Incorrect punctuation	Correctness
104.	they likely use	Wordy sentences	Clarity
105.	*is not explicitly mentioned*	Passive voice misuse	Clarity
106.	*is assigned*	Passive voice misuse	Clarity
107.	~~of~~ → to	Wrong or missing prepositions	Correctness
108.	*is recomputed*	Passive voice misuse	Clarity
109.	a confidence	Determiner use (a/an/the/this, etc.)	Correctness

110.	*is considered*	Passive voice misuse	Clarity
111.	~~gaussian~~ → A Gaussian	Determiner use (a/an/the/this, etc.)	Correctness
112.	*is applied*	Passive voice misuse	Clarity
113.	~~gaussian~~ → Gaussian	Confused words	Correctness
114.	*is expressed*	Passive voice misuse	Clarity
115.	*are then downsampled*	Passive voice misuse	Clarity
116.	~~resolu~~ → resolution	Misspelled words	Correctness
117.	~~so as to~~ → to	Wordy sentences	Clarity
118.	~~gaussian~~ → Gaussian	Confused words	Correctness
119.	~~applied~~	Incorrect verb forms	Correctness
120.	*is applied*	Passive voice misuse	Clarity
121.	~~average pooled~~ → average-pooled	Misspelled words	Correctness
122.	~~max pooled~~ → max-pooled	Misspelled words	Correctness
123.	*is taken*	Passive voice misuse	Clarity
124.	~~and is~~ → . It is	Hard-to-read text	Clarity
125.	*is then fed*	Passive voice misuse	Clarity
126.	, which	Incorrect punctuation	Correctness
127.	*is passed*	Passive voice misuse	Clarity
128.	*is passed on*	Passive voice misuse	Clarity

129.	, which	Incorrect punctuation	Correctness
130.	*is trained*	Passive voice misuse	Clarity
131.	~~cross entropy~~ → cross-entropy	Misspelled words	Correctness
132.	, and	Incorrect punctuation	Correctness
133.	~~told~~ → said	Incorrect verb forms	Correctness
134.	*The ambiguity arises on how this operation is exactly done?*	Ungrammatical sentence	Correctness
135.	*is exactly done*	Passive voice misuse	Clarity
136.	The first interpretation is likely	Wordy sentences	Clarity
137.	*is used*	Passive voice misuse	Clarity
138.	as good	Misuse of modifiers	Correctness
139.	~~to~~	Incorrect verb forms	Correctness
140.	~~OADN,~~ → OADN;	Incorrect punctuation	Correctness
141.	*are trained*	Passive voice misuse	Clarity
142.	, where	Incorrect punctuation	Correctness
143.	The family	Determiner use (a/an/the/this, etc.)	Correctness
144.	~~on~~ → of	Wrong or missing prepositions	Correctness
145.	*is shown*	Passive voice misuse	Clarity
146.	~~figure~~ → Figure	Confused words	Correctness
147.	*are preprocessed*	Passive voice misuse	Clarity

148.	*is cropped*	Passive voice misuse	Clarity
149.	~~be of the~~ → a	Incorrect phrasing	Correctness
150.	, and	Incorrect punctuation	Correctness
151.	*are normalized*	Passive voice misuse	Clarity
152.	~~normalized~~ → normalised	Mixed dialects of English	Correctness
153.	] .	Improper formatting	Correctness
154.	, which	Incorrect punctuation	Correctness
155.	~~discriminator~~ → discriminators	Incorrect noun number	Correctness
156.	, and	Incorrect punctuation	Correctness
157.	a reconstructed	Determiner use (a/an/the/this, etc.)	Correctness
158.	*is trained*	Passive voice misuse	Clarity
159.	~~kernels~~ → kernel	Incorrect noun number	Correctness
160.	a stride	Determiner use (a/an/the/this, etc.)	Correctness
161.	of 1	Wrong or missing prepositions	Correctness
162.	*is followed*	Passive voice misuse	Clarity
163.	~~normalization~~ → normalisation	Mixed dialects of English	Correctness
164.	decoder's output	Paragraph can be improved	Clarity
165.	~~as~~ → to	Wrong or missing prepositions	Correctness

166.	, but	Incorrect punctuation	Correctness
167.	] .	Improper formatting	Correctness
168.	with a stride	Wrong or missing prepositions	Correctness
169.	*is not mentioned*	Passive voice misuse	Clarity
170.	*is used*	Passive voice misuse	Clarity
171.	it is likely	Wordy sentences	Clarity
172.	, considering	Incorrect punctuation	Correctness
173.	~~uses~~ → use	Faulty subject-verb agreement	Correctness
174.	a stride	Determiner use (a/an/the/this, etc.)	Correctness
175.	of 2	Wrong or missing prepositions	Correctness
176.	~~the~~ models	Determiner use (a/an/the/this, etc.)	Correctness
177.	with a stride	Wrong or missing prepositions	Correctness
178.	*is not mentioned*	Passive voice misuse	Clarity
179.	*is used*	Passive voice misuse	Clarity
180.	it is likely	Wordy sentences	Clarity
181.	, considering	Incorrect punctuation	Correctness
182.	~~uses~~ → use	Faulty subject-verb agreement	Correctness

183.	a stride	Determiner use (a/an/the/this, etc.)	Correctness
184.	of 2	Wrong or missing prepositions	Correctness
185.	~~the~~ models	Determiner use (a/an/the/this, etc.)	Correctness
186.	, coming	Incorrect punctuation	Correctness
187.	~~which tells~~ → that tells	Pronoun use	Correctness
188.	which,	Incorrect punctuation	Correctness
189.	~~lay~~ → , lies	Incorrect phrasing	Correctness
190.	~~ers~~ → The model	Incorrect phrasing	Correctness
191.	*are trained*	Passive voice misuse	Clarity
192.	~~which do~~ → that do	Pronoun use	Correctness
193.	*This*	Intricate text	Clarity
194.	*is done*	Passive voice misuse	Clarity
195.	*is trained*	Passive voice misuse	Clarity
196.	the weighted	Determiner use (a/an/the/this, etc.)	Correctness
197.	*This*	Intricate text	Clarity
198.	pixel-level similarity	Incorrect phrasing	Correctness
199.	~~larity~~ → Similarity	Misspelled words	Correctness
200.	, which	Incorrect punctuation	Correctness

201.	, and	Incorrect punctuation	Correctness
202.	~~which~~ → Which	Confused words	Correctness
203.	~~area.~~ → area?	Closing punctuation	Correctness
204.	*is introduced*	Passive voice misuse	Clarity
205.	d1 .	Improper formatting	Correctness
206.	the adversarial	Determiner use (a/an/the/this, etc.)	Correctness
207.	, which	Incorrect punctuation	Correctness
208.	~~cross entropy~~ → cross-entropy	Misspelled words	Correctness
209.	*are trained*	Passive voice misuse	Clarity
210.	alternatively .	Improper formatting	Correctness
211.	*is trained*	Passive voice misuse	Clarity
212.	*is given*	Passive voice misuse	Clarity
213.	~~figure~~ → Figure	Confused words	Correctness
214.	*are derived*	Passive voice misuse	Clarity
215.	a histogram	Determiner use (a/an/the/this, etc.)	Correctness
216.	*is defined*	Passive voice misuse	Clarity
217.	, and	Incorrect punctuation	Correctness
218.	~~th of~~	Incorrect phrasing	Correctness
219.	~~autoen~~ → autoencoder	Incorrect noun number	Correctness

220.	*coder which assigns a confidence score to each triangle.*	Ungrammatical sentence	Correctness
221.	, however	Incorrect punctuation	Correctness
222.	*was also shown*	Passive voice misuse	Clarity
223.	*is given*	Passive voice misuse	Clarity
224.	~~figure~~ → Figure	Confused words	Correctness
225.	~~deal~~ → deals	Faulty subject-verb agreement	Correctness
226.	~~this is~~ → which is	Pronoun use	Correctness
227.	, and	Incorrect punctuation	Correctness
228.	~~cross entropy~~ → cross-entropy	Misspelled words	Correctness
229.	*are fixed*	Passive voice misuse	Clarity
230.	*is fine tuned*	Passive voice misuse	Clarity
231.	~~fine tuned~~ → fine-tuned	Misspelled words	Correctness
232.	~~the output~~ → The outputs	Incorrect phrasing	Correctness
233.	, and	Incorrect punctuation	Correctness
234.	a loss	Determiner use (a/an/the/this, etc.)	Correctness
235.	~~regularization~~ → regularisation	Mixed dialects of English	Correctness
236.	is that	Incorrect verb forms	Correctness
237.	A privileged	Determiner use (a/an/the/this, etc.)	Correctness

238.	The classifier	Determiner use (a/an/the/this, etc.)	Correctness
239.	but it	Conjunction use	Correctness
240.	*is not used*	Passive voice misuse	Clarity
241.	*is shown*	Passive voice misuse	Clarity
242.	~~figure~~ → Figure	Confused words	Correctness
243.	~~gabor~~ → Gabor	Confused words	Correctness
244.	, giving	Incorrect punctuation	Correctness
245.	~~gabor~~ → Gabor	Confused words	Correctness
246.	*are extracted*	Passive voice misuse	Clarity
247.	, and	Incorrect punctuation	Correctness
248.	~~chosen,~~ → chosen;	Incorrect punctuation	Correctness
249.	~~gabor~~ → Gabor	Confused words	Correctness
250.	*This*	Intricate text	Clarity
251.	*is done*	Passive voice misuse	Clarity
252.	~~to~~ → for	Wrong or missing prepositions	Correctness
253.	*is introduced*	Passive voice misuse	Clarity
254.	~~matching,~~ → matching.	Incorrect punctuation	Correctness
255.	~~centered~~ → , centred	Incorrect punctuation	Correctness
256.	*is slid*	Passive voice misuse	Clarity

257.	28000-dimensional	Misspelled words	Correctness
258.	*is trained*	Passive voice misuse	Clarity
259.	, and	Incorrect punctuation	Correctness
260.	*is calculated*	Passive voice misuse	Clarity
261.	, and	Incorrect punctuation	Correctness
262.	the templates	Determiner use (a/an/the/this, etc.)	Correctness
263.	*are removed*	Passive voice misuse	Clarity
264.	*This*	Intricate text	Clarity
265.	*is done*	Passive voice misuse	Clarity
266.	*are reduced*	Passive voice misuse	Clarity
267.	*are reduced*	Passive voice misuse	Clarity
268.	~~to~~ → and	Conjunction use	Correctness
269.	the final	Determiner use (a/an/the/this, etc.)	Correctness
270.	~~figure~~ → Figure	Confused words	Correctness
271.	*The framework of RGBT is given in figure 7.*	Paragraph can be improved	Clarity
272.	~~dual branch~~ → dual-branch	Misspelled words	Correctness
273.	, which	Incorrect punctuation	Correctness
274.	*is applied*	Passive voice misuse	Clarity
275.	training,	Incorrect punctuation	Correctness

276.	*is randomly flipped*	Passive voice misuse	Clarity
277.	, and	Incorrect punctuation	Correctness
278.	~~color~~ → colour	Mixed dialects of English	Correctness
279.	the local	Determiner use (a/an/the/this, etc.)	Correctness
280.	the first	Determiner use (a/an/the/this, etc.)	Correctness
281.	, and	Incorrect punctuation	Correctness
282.	*is fed*	Passive voice misuse	Clarity
283.	*This*	Intricate text	Clarity
284.		Tone suggestions	Delivery
285.	*This*	Intricate text	Clarity
286.	*is not used*	Passive voice misuse	Clarity
287.	~~classification,~~ → classification;	Incorrect punctuation	Correctness
288.	~~Transformer~~ → The transformer	Determiner use (a/an/the/this, etc.)	Correctness
289.	of which are	Pronoun use	Correctness
290.	*are processed*	Passive voice misuse	Clarity
291.	a regular	Determiner use (a/an/the/this, etc.)	Correctness
292.	*are computed*	Passive voice misuse	Clarity
293.	*are then used*	Passive voice misuse	Clarity

294.	*is used*	Passive voice misuse	Clarity
295.	~~the~~ branches	Determiner use (a/an/the/this, etc.)	Correctness
296.	*are passed*	Passive voice misuse	Clarity
297.	*nected layer to get P red2 which is our final prediction [27].*	Ungrammatical sentence	Correctness
298.	*is given*	Passive voice misuse	Clarity
299.	~~figure~~ → Figure	Confused words	Correctness
300.	*This*	Intricate text	Clarity
301.	*these models could not be tested*	Passive voice misuse	Clarity
302.	*must be noted*	Passive voice misuse	Clarity
303.	~~originally~~	Wordy sentences	Clarity
304.	*almost all of the above datasets were originally created*	Passive voice misuse	Clarity
305.	result,	Incorrect punctuation	Correctness
306.	*was collected*	Passive voice misuse	Clarity
307.	~~a~~ → an	Determiner use (a/an/the/this, etc.)	Correctness
308.	*This*	Intricate text	Clarity
309.	*was collected*	Passive voice misuse	Clarity
310.	*were removed*	Passive voice misuse	Clarity
311.	*This*	Intricate text	Clarity

312.	, too	Incorrect punctuation	Correctness
313.	, is	Incorrect punctuation	Correctness
314.	~~real world~~ → real-world	Misspelled words	Correctness
315.	*This*	Intricate text	Clarity
316.	a large	Determiner use (a/an/the/this, etc.)	Correctness
317.	~~who~~ → that were	Incorrect phrasing	Correctness
318.	*This*	Intricate text	Clarity
319.	, as	Incorrect punctuation	Correctness
320.	~~, but using~~ → . Still, using	Hard-to-read text	Clarity
321.	~~models~~ → models'	Incorrect noun number	Correctness
322.	*This*	Intricate text	Clarity
323.	*This is because there is no single figure provided by the authors.*	Passive voice misuse	Clarity
324.	*most of the above models were trained*	Passive voice misuse	Clarity
325.	~~times~~	Incorrect noun number	Correctness
326.	~~times~~	Wordy sentences	Clarity
327.	~~Both~~ → Neither	Determiner use (a/an/the/this, etc.)	Correctness
328.	~~were not~~ → was	Faulty subject-verb agreement	Correctness
329.	*were not collected*	Passive voice misuse	Clarity

330.	*are tested*	Passive voice misuse	Clarity
331.	Figure 4	Incorrect noun number	Correctness
332.	*were tested*	Passive voice misuse	Clarity
333.	occlusion-specific	Misspelled words	Correctness
334.	, which	Incorrect punctuation	Correctness
335.	model,	Incorrect punctuation	Correctness
336.	~~patch based~~ → patch-based	Misspelled words	Correctness
337.	*This*	Intricate text	Clarity
338.	~~real world~~ → real-world	Misspelled words	Correctness
339.	, which	Incorrect punctuation	Correctness
340.	, and	Incorrect punctuation	Correctness
341.	*nothing can be said*	Passive voice misuse	Clarity
342.	*it can be explained*	Passive voice misuse	Clarity
343.	~~realizing~~ → realising	Mixed dialects of English	Correctness
344.		Tone suggestions	Delivery
345.	~~ar~~ → far	Misspelled words	Correctness
346.	, but	Incorrect punctuation	Correctness
347.		Tone suggestions	Delivery
348.	~~marks~~ → Marks	Confused words	Correctness
349.	~~choose~~ → chose	Incorrect verb forms	Correctness

350.	*This*	Intricate text	Clarity
351.	~~has~~ → have	Faulty subject-verb agreement	Correctness
352.	~~on~~	Wrong or missing prepositions	Correctness
353.	*This*	Intricate text	Clarity
354.	real-world conditions	Incorrect phrasing	Correctness
355.	~~and~~	Conjunction use	Correctness
356.	*a large, high quality facial images datasets that has occlusions.*	Ungrammatical sentence	Correctness
357.	*a large, high quality facial images datasets that has occlusions.*	Paragraph can be improved	Clarity
358.	*This*	Intricate text	Clarity
359.	~~prior~~ → previously	Misuse of modifiers	Correctness
360.	~~occuring~~ → occurring	Misspelled words	Correctness
361.	~~real world~~ → real-world	Misspelled words	Correctness
362.	occlusion-specific	Misspelled words	Correctness
363.	features,	Incorrect punctuation	Correctness
364.	*This*	Intricate text	Clarity
365.	~~the FED-RO~~ → The FED-RO	Confused words	Correctness
366.	~~real world~~ → real-world	Misspelled words	Correctness
367.	*models using privileged information to deal with OFER.*	Paragraph can be improved	Clarity

368.	*This*	Intricate text	Clarity
369.	The CK+	Determiner use (a/an/the/this, etc.)	Correctness
370.	~~multiple~~ → Multiple	Confused words	Correctness
371.	*the creation of a high accuracy OFER model.*	Ungrammatical sentence	Correctness
372.	~~when~~ → When	Confused words	Correctness
373.	~~by adding some~~ → Adding an	Incorrect phrasing	Correctness
374.	~~of~~ → as	Wrong or missing prepositions	Correctness
375.	~~give~~ → Give	Confused words	Correctness